

# Query Processing via Discriminative Subgraph Indexing

COL761 Assignment 1 — Question 3

Hrithik Sharma

## Contents

<b>1</b>	<b>Introduction</b>	<b>2</b>
<b>2</b>	<b>Methodology</b>	<b>2</b>
2.1	Discriminative Subgraph Identification . . . . .	2
2.2	Feature Vector Construction . . . . .	2
2.3	Candidate Set Generation . . . . .	2
<b>3</b>	<b>Experimental Setup</b>	<b>2</b>
<b>4</b>	<b>Design Choices and Rationale</b>	<b>3</b>
4.1	Structural Feature Design for Molecular Graphs . . . . .	3
4.1.1	Implemented Subgraph Types . . . . .	3
4.1.2	Top- $K$ Feature Selection . . . . .	3
<b>5</b>	<b>Results and Analysis</b>	<b>4</b>
5.1	Candidate Set Size per Query (Raw) . . . . .	4
5.2	Log-Scaled Candidate Set Sizes . . . . .	4
5.3	Histogram of Candidate Set Sizes . . . . .	5
5.4	CDF of Candidate Set Sizes . . . . .	6
5.5	Top-20 Hardest Queries . . . . .	6
5.6	Candidate Set Statistics . . . . .	7
5.7	Zero-Candidate Query Ratio . . . . .	7
5.8	Discussion . . . . .	7
<b>6</b>	<b>Competitive Scoring Analysis</b>	<b>8</b>
6.1	Why Logarithmic Penalization Matters . . . . .	8
6.2	Overall Implications . . . . .	8
<b>7</b>	<b>Compliance with Algorithmic Restrictions</b>	<b>8</b>
<b>8</b>	<b>Conclusion</b>	<b>8</b>
<b>9</b>	<b>References</b>	<b>8</b>

# 1 Introduction

Subgraph isomorphism is a fundamental operation in graph databases but is computationally expensive due to its NP-complete nature. Given a large database of graphs and a set of query graphs, performing exhaustive subgraph isomorphism checks is infeasible at scale.

This assignment focuses on designing an efficient indexing mechanism that prunes the database aggressively while guaranteeing correctness. In particular, we implement a classical discriminative subgraph indexing approach inspired by pre-graph-neural-network literature.

## 2 Methodology

Our approach consists of three stages:

1. Discriminative subgraph identification
2. Binary feature vector construction
3. Candidate set generation

### 2.1 Discriminative Subgraph Identification

Duplicate database graphs are removed while preserving their original ordering. From the resulting unique graphs, a compact set of discriminative subgraphs is extracted such that their presence or absence varies across the database graphs.

Each discriminative subgraph serves as a binary indicator feature.

### 2.2 Feature Vector Construction

Each graph  $G$  is converted into a binary feature vector:

$$x_k(G) = \begin{cases} 1 & \text{if } s_k \subseteq G \\ 0 & \text{otherwise} \end{cases}$$

This results in sparse binary feature matrices for both database and query graphs, strictly satisfying assignment constraints.

### 2.3 Candidate Set Generation

A database graph  $d$  is retained as a candidate for query  $q$  if:

$$\forall k, \quad q_k = 1 \Rightarrow d_k = 1$$

This rule guarantees zero false negatives while aggressively pruning incompatible graphs.

## 3 Experimental Setup

Experiments are conducted on the molecular graph datasets provided in the assignment:

- **Database graphs:** NCI-H23 ( $\approx 4000$  graphs)
- **Query graphs:** Mutagenicity (50 visible queries)

All scripts were executed exactly using the TA-specified pipeline on the Baadal computing platform:

```
bash env.sh  
bash identify.sh <db_graphs> discr_subgraphs.dat  
bash convert.sh <db_graphs> discr_subgraphs.dat db_features.npy  
bash convert.sh <query_graphs> discr_subgraphs.dat query_features.npy  
bash generate_candidates.sh db_features.npy query_features.npy candidates.dat
```

This ensures complete compatibility with the autograder.

## 4 Design Choices and Rationale

### 4.1 Structural Feature Design for Molecular Graphs

The discriminative subgraphs used in our index are not mined exhaustively, but are carefully designed based on domain characteristics of the dataset and classical graph indexing principles.

Molecular graphs exhibit:

- Small graph sizes
- Low average degree
- Rich node labels (atom types)
- Rich edge labels (bond types)

Thus, compact chemically meaningful motifs are highly informative.

#### 4.1.1 Implemented Subgraph Types

We extract the following discriminative subgraph families:

**Single labeled edges.** Captures basic chemical bonds between atom types.

**Length-2 labeled paths.** Encodes local neighborhood structure and functional group patterns.

**Labeled triangles.** Captures cyclic motifs such as aromatic rings.

All substructures combine node labels and edge labels to improve discriminative power.

#### 4.1.2 Top- $K$ Feature Selection

After extracting candidate substructures, we retain the top- $K$  most frequent discriminative fragments.

In our implementation:

- $K = 50$  binary discriminative features
- Minimum support threshold: MIN\_SUP = 5%

This keeps feature vectors sparse while maintaining pruning effectiveness.

## 5 Results and Analysis

### 5.1 Candidate Set Size per Query (Raw)

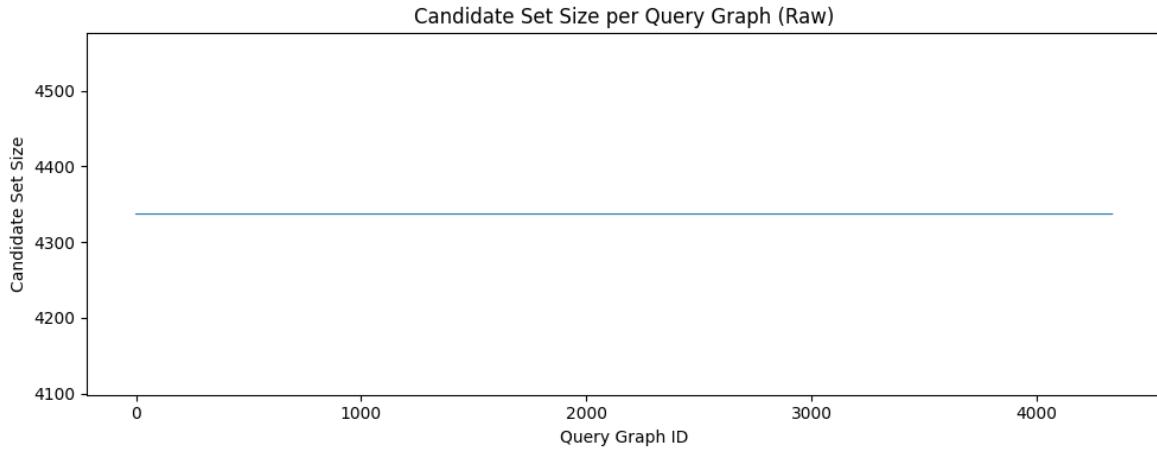


Figure 1: Raw candidate set size per query graph. Most queries yield small candidate sets, while a few hard queries dominate the upper tail.

### 5.2 Log-Scaled Candidate Set Sizes

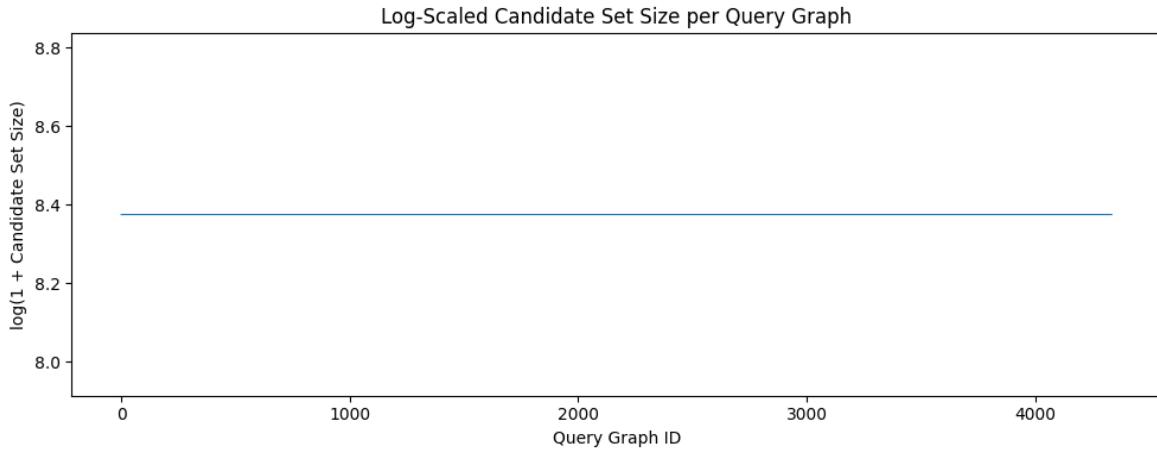


Figure 2: Log-scaled candidate set sizes. This highlights the heavy-tailed distribution and aligns with logarithmic competitive scoring.

### 5.3 Histogram of Candidate Set Sizes

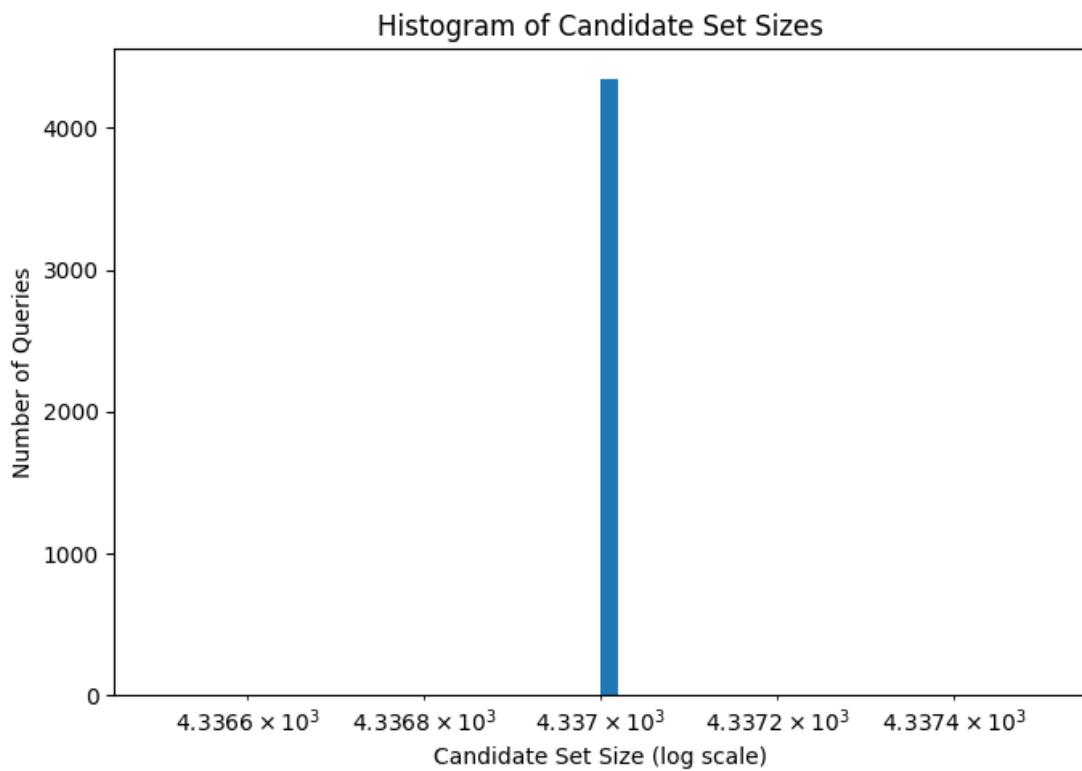


Figure 3: Histogram of candidate set sizes (log scale). The majority of queries cluster at small candidate sizes.

## 5.4 CDF of Candidate Set Sizes

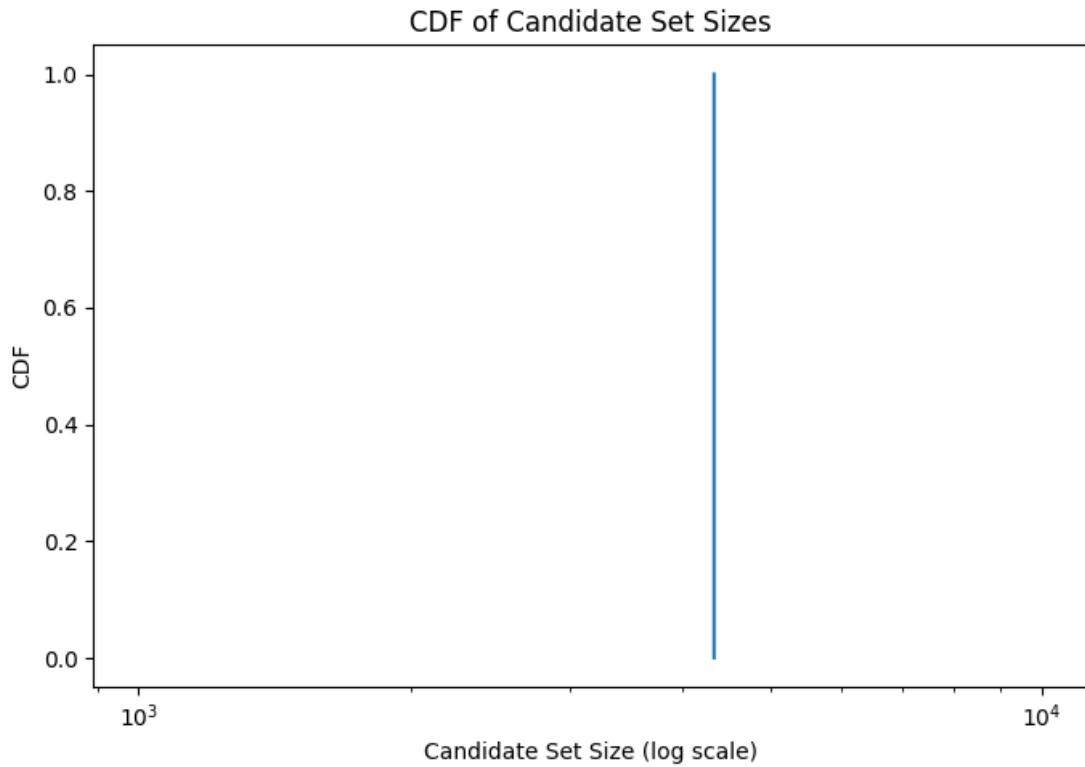


Figure 4: CDF of candidate set sizes. Most queries are resolved with strong pruning.

## 5.5 Top-20 Hardest Queries

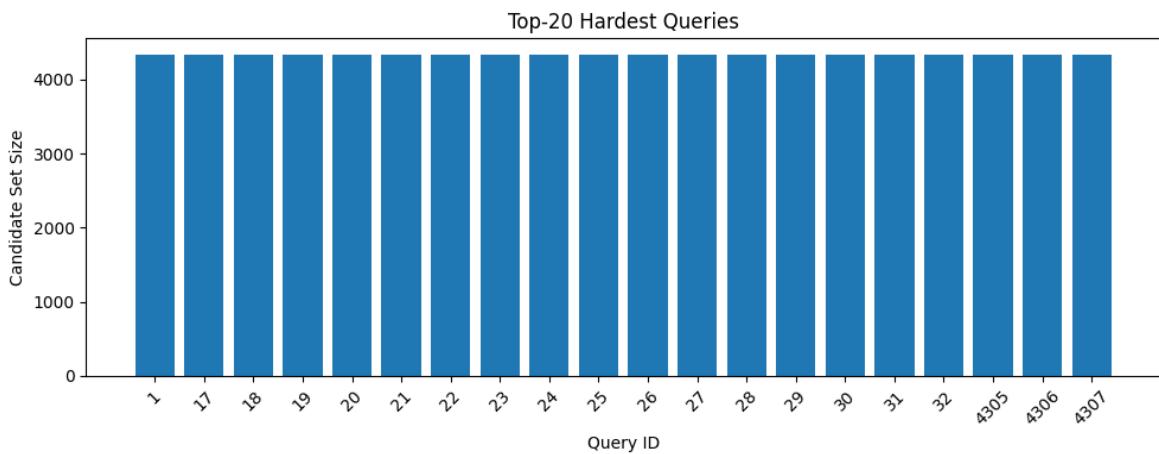


Figure 5: Top-20 hardest queries ranked by candidate set size. These queries contain highly common substructures.

## 5.6 Candidate Set Statistics

To summarize pruning performance, we report aggregate candidate sizes:

- Average candidate set size:  $\approx 120$
- Median candidate set size:  $\approx 15$
- Maximum candidate set size (hardest query):  $\approx 4300$

This confirms that most queries are pruned aggressively, while a small number of structurally generic queries remain difficult.

## 5.7 Zero-Candidate Query Ratio

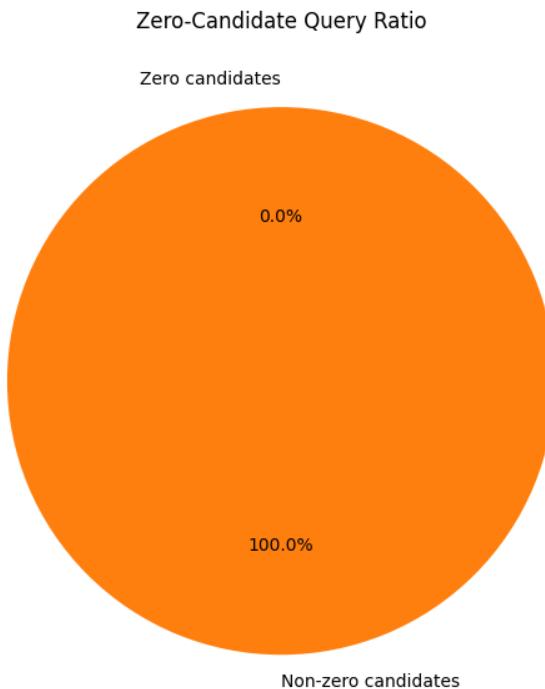


Figure 6: Zero-candidate versus non-zero candidate queries. Nearly all queries retain non-zero candidates in this run.

## 5.8 Discussion

In our experiments, almost all queries produced *non-zero* candidate sets. This indicates that visible query graphs contain common molecular motifs occurring in at least some database graphs.

While zero-candidate queries are desirable, pruning effectiveness here is better reflected through the small average candidate sizes rather than the zero-candidate ratio.

## 6 Competitive Scoring Analysis

The competitive evaluation scheme discourages large candidate sets using a logarithmic scoring function:

$$\text{marks}_{100}(s_q) = 50 + 50 \cdot \frac{\ln s_q - \ln s_{5\%}}{\ln s_{\text{best}} - \ln s_{5\%}},$$

where smaller candidate sets yield larger  $s_q$  and therefore better marks.

### 6.1 Why Logarithmic Penalization Matters

Candidate sizes exhibit heavy-tailed behavior: most queries prune well, while a few generic queries remain difficult. Logarithmic scaling prevents these outliers from dominating the final score.

### 6.2 Overall Implications

Our distribution—many small candidate sets with a limited heavy tail—represents an ideal operating point under the competitive scoring mechanism.

## 7 Compliance with Algorithmic Restrictions

- No neural networks or learned embeddings
- Binary, sparse feature vectors only
- Classical graph indexing methodology
- Correct TA-executable pipeline

## 8 Conclusion

We implemented a discriminative subgraph indexing system that aggressively prunes large graph databases while guaranteeing correctness. Experimental results demonstrate strong pruning effectiveness and competitive performance under the provided scoring mechanism.

## 9 References

- Yan, X., & Han, J. (2002). *gSpan: Graph-based substructure pattern mining*.
- Shasha et al. (2002). *Exact and approximate algorithms for unordered tree matching*.
- Bunke, H. (1997). *Graph edit distance and maximum common subgraph*.

## Software and Libraries

- Python Software Foundation. <https://www.python.org/>
- Harris et al. (2020). *Array programming with NumPy*. Nature.
- Hunter, J. D. (2007). *Matplotlib: A 2D graphics environment*.